

Deep Learning Image Processing

# Case Study: LeNet-5 & AlexNet

---

2023-1

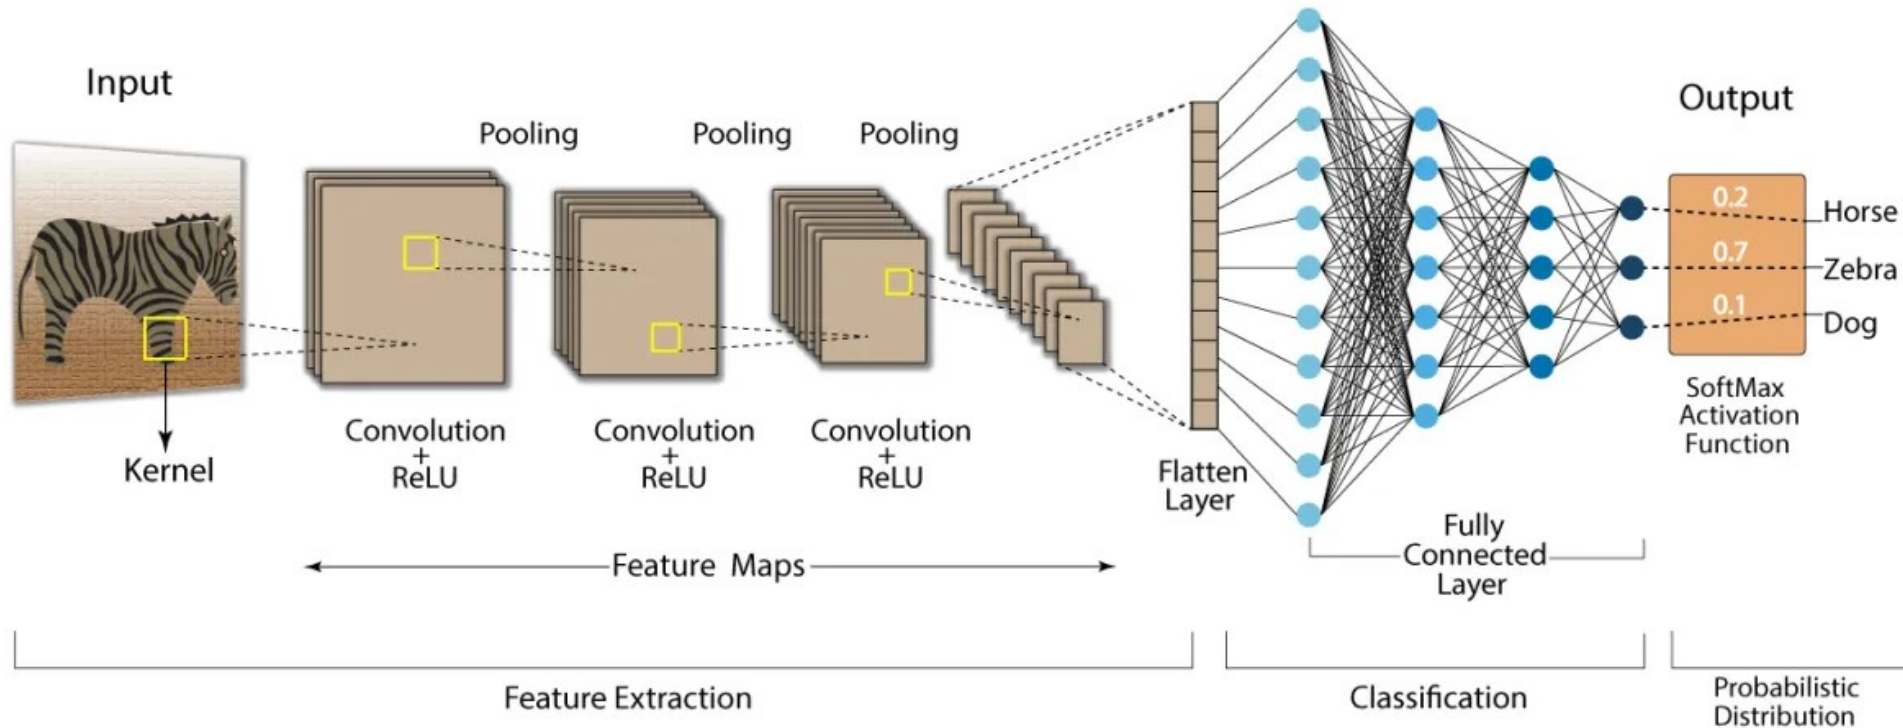
Young-Keun Kim

School of Mechanical & Control Engineering

# Overview

## ● CNN architecture

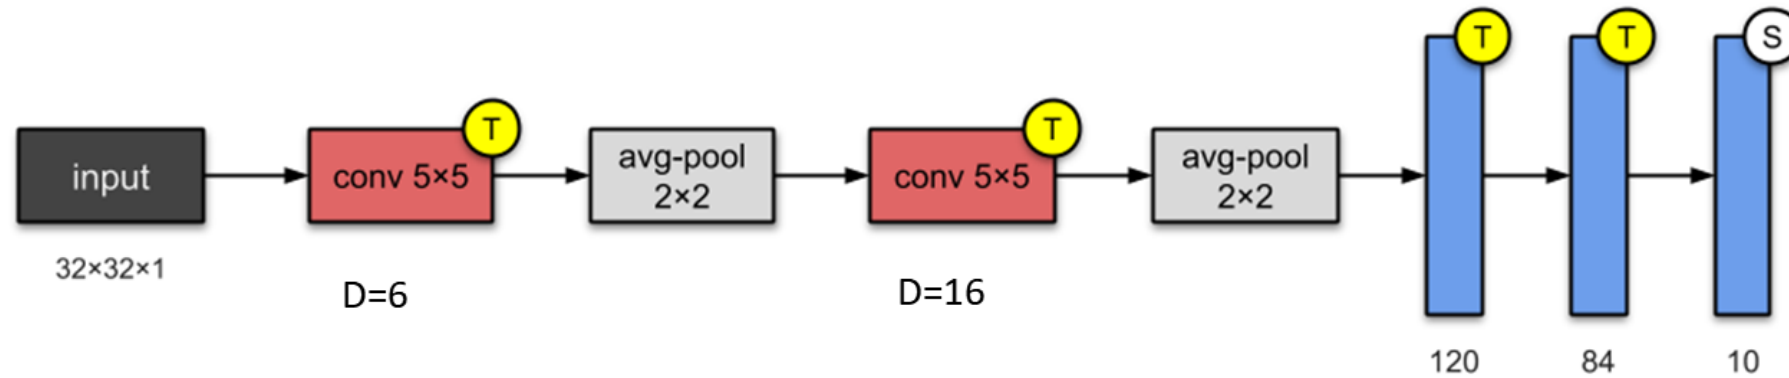
- For Classification Prediction of an image
  - Feature Extraction - Classification - Probability Distribution
- Basic CNN architecture components
  - convolutional layer, activation function, pooling layer, fully-connected layer



# LeNet-5

## ● Case study of LeNet-5

- LeCun, Yann et al., Gradient-based learning applied to document recognition, 1998
  - For handwriting recognition, MNIST
- A standard template of CNN
- 2 CONV layers + 3 F.C layers
- About 60,000 parameters

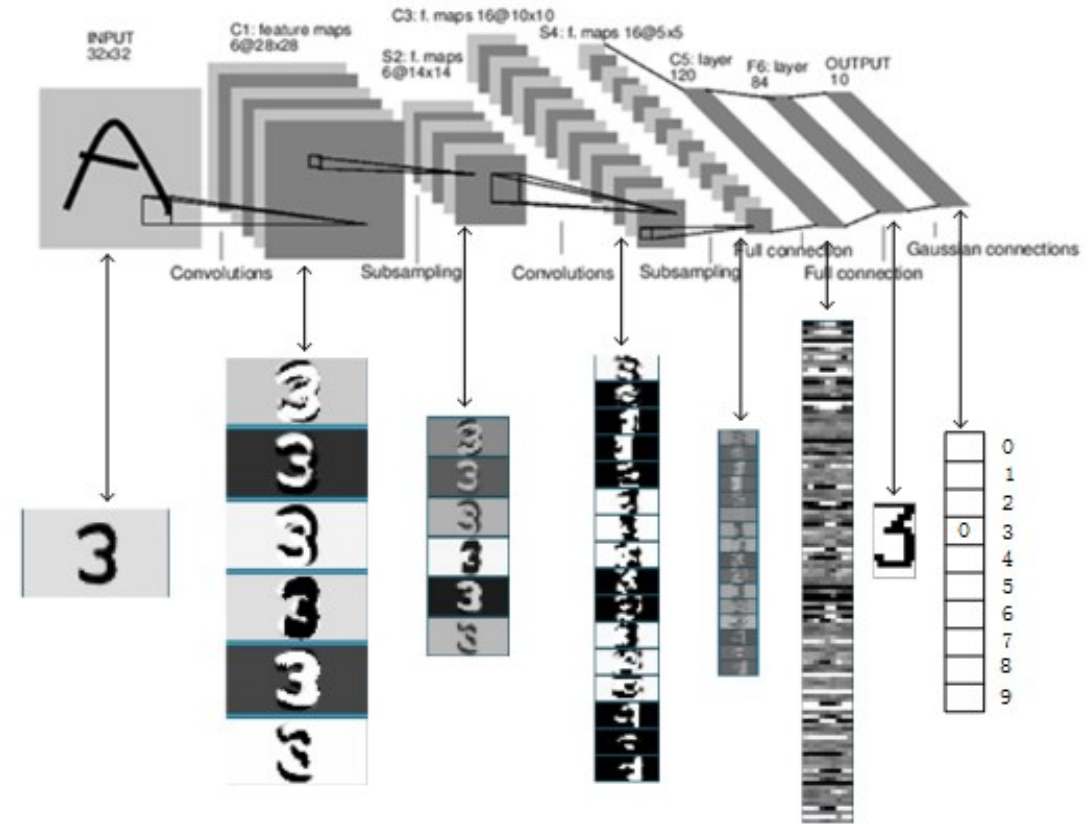


[CNN architecture of LeNet-5]

# LeNet-5

## ● Case study of LeNet-5

- A standard template of CNN
- Activation Function: TanH
  - Now, ReLU is often used
- Pooling: Avg. pooling
  - Now, MaxPooling is often used
- No Padding
- F.C: softmax
  - Originally used RBF(Radial Basis Function)
- Loss Function: MSE
- Input: 32x32x1
  - MNIST image is 28x28
  - MNIST is padded to 32



# Case study : LeNet-5

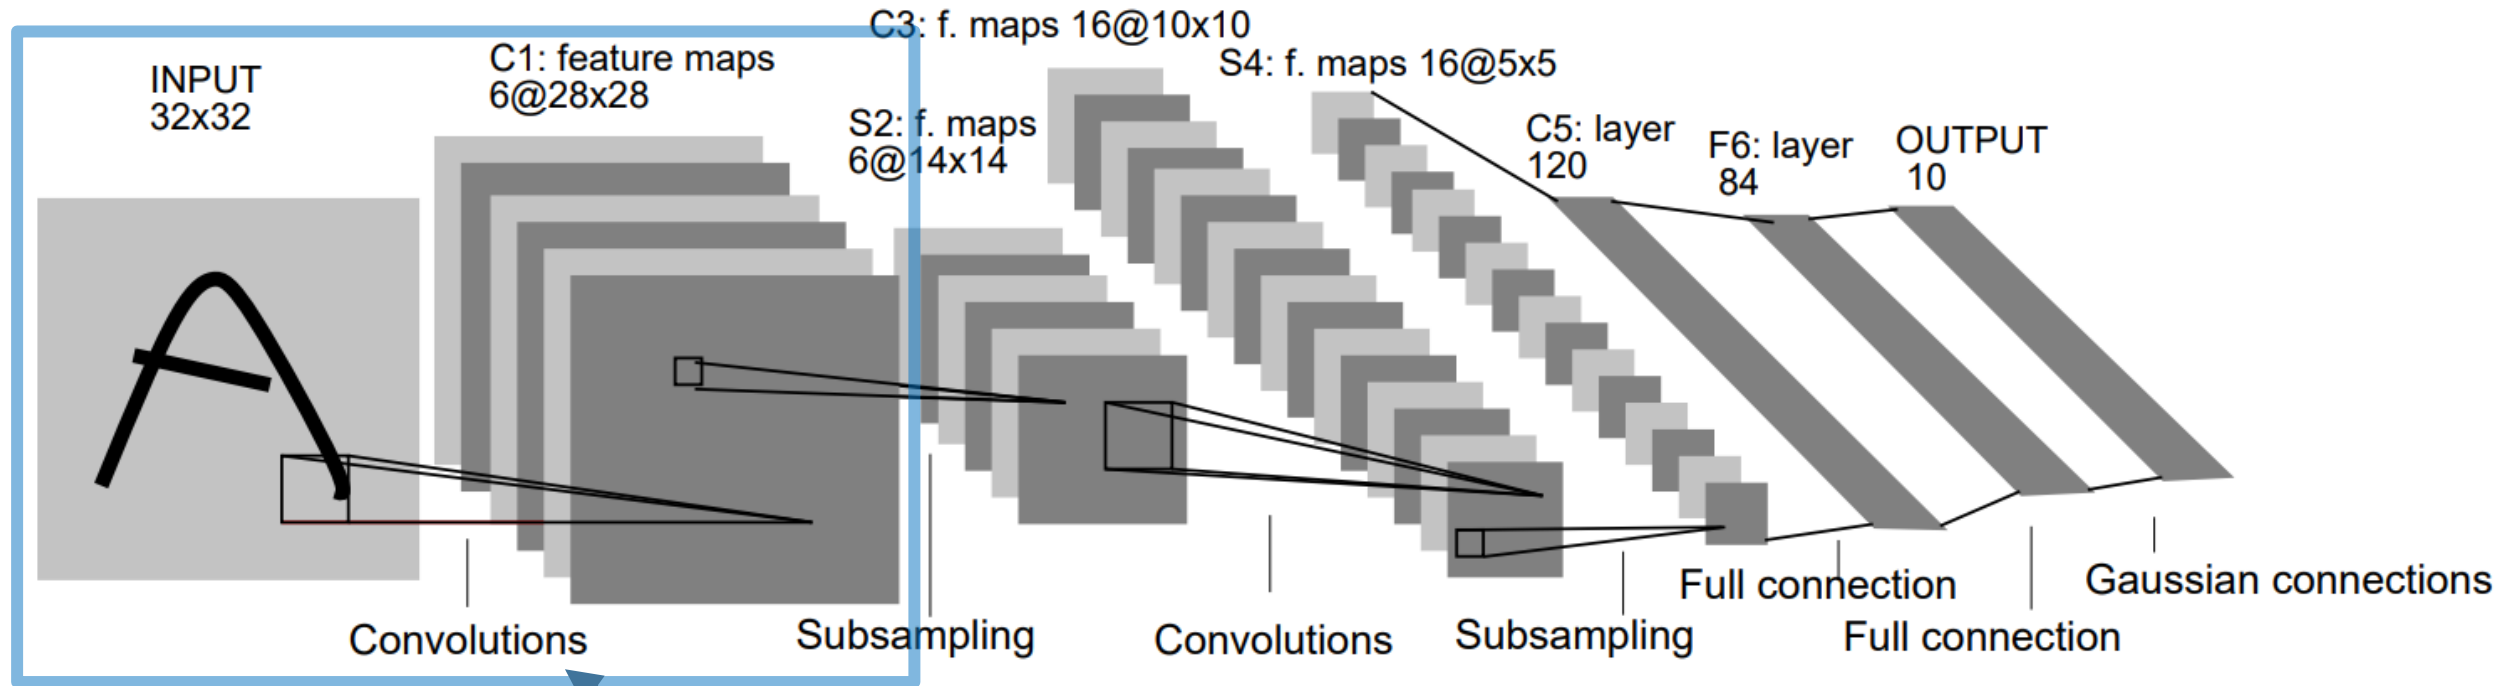


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

without zero-padding & stride=1

If zero-padding-> $(5*5+1)*32*32*6$

Step	No. of free parameters	No. of connections
Input→C1 : 6 of convolution filter kernel (5x5x1)	$156=(5*5+1(\text{bias}))*6$	$122,304=(5*5+1)*28*28*6$

\*MLP:  $(32*32+1)*28*28*6=4,821,600$

\*MLP:  $(32*32+1)*28*28*6=4,821,600$

# Case study : LeNet-5

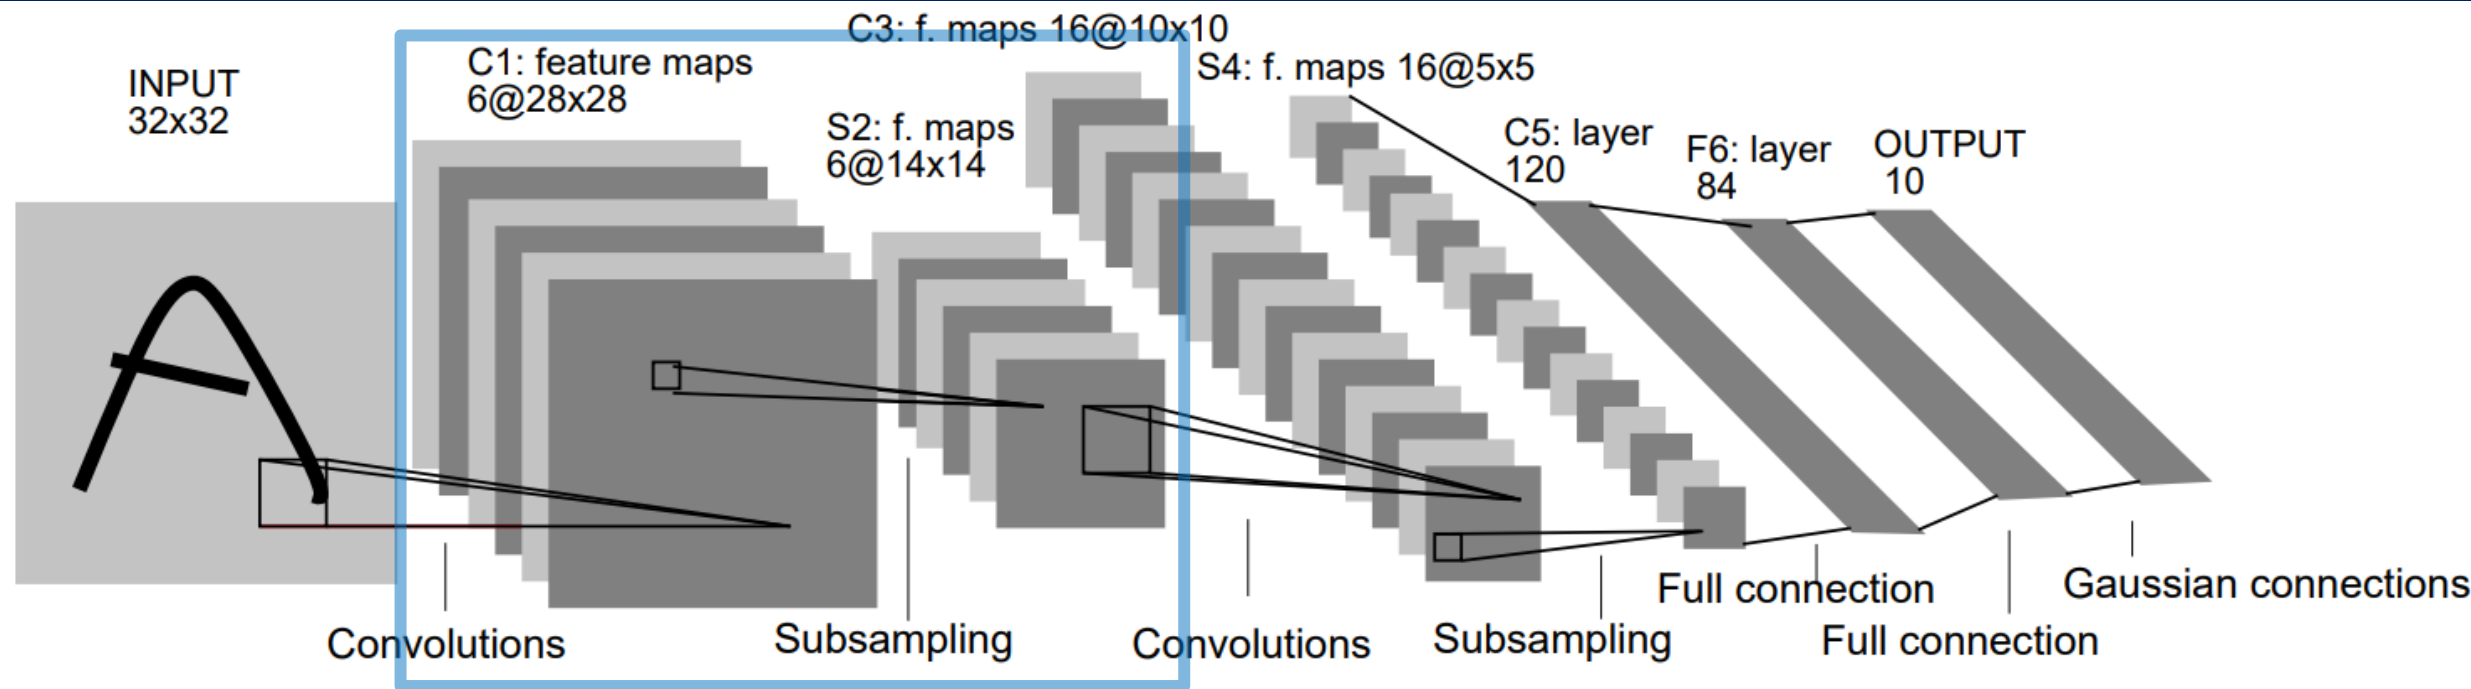


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Step	No. of free parameters	No. of connections
C1→S2 : average pooling (2x2), stride 2	$12 = (1(\text{weight}) + 1(\text{bias})) * 6$ (* usually, pooling has no weight)	$5880 = (2 * 2 + 1) * 14 * 14 * 6$

# Case study : LeNet-5




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Step	No. of free parameters	No. of connections
C2→S3 : convolution kernel (5x5)	1516=5*5*60+16	151600=(5*5*60+16)*10*10

# Case study : LeNet-5

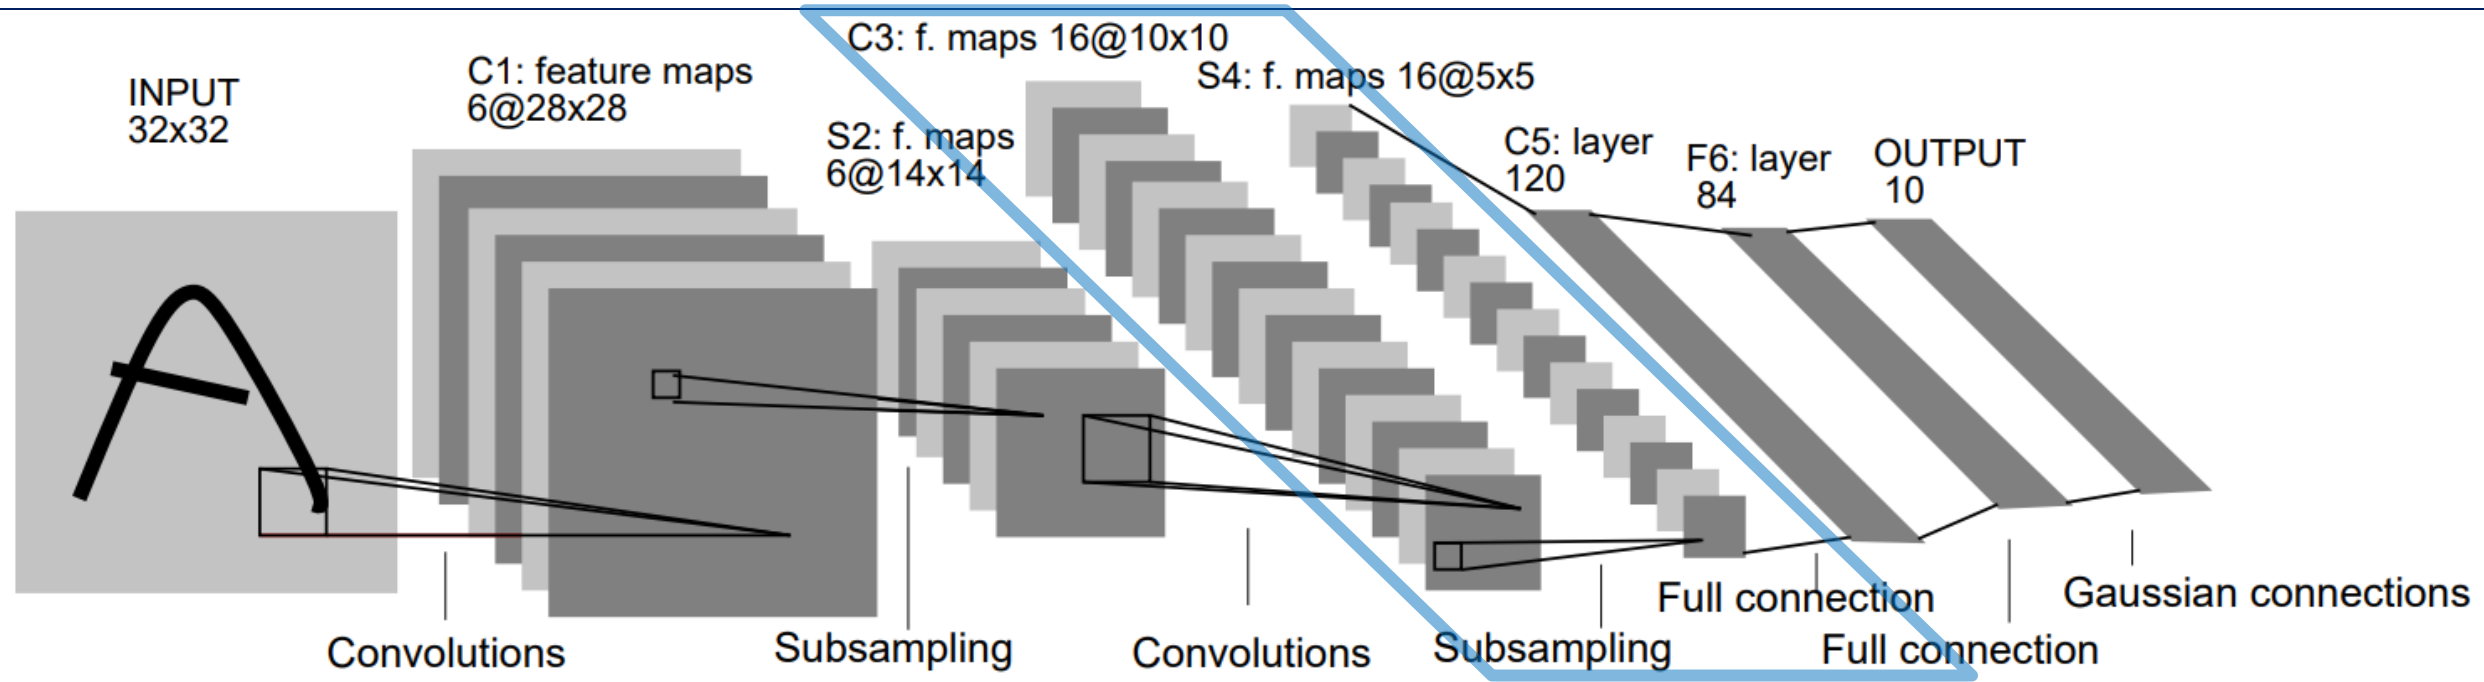


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Step	No. of free parameters	No. of connections
C3→S4 : average pooling (2x2)	$32=(1(\text{weight})+1(\text{bias})) \cdot 16$	$2000=(2 \cdot 2+1) \cdot 5 \cdot 5 \cdot 16$



# Case study : LeNet-5




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Step	No. of free parameters	No. of connections
S4→C5 : convolution kernel (5x5)	$48,120 = 5 * 5 * 16 * 120 + 120(\text{bias})$	Same as 'No. of free parameters'

# Case study : LeNet-5

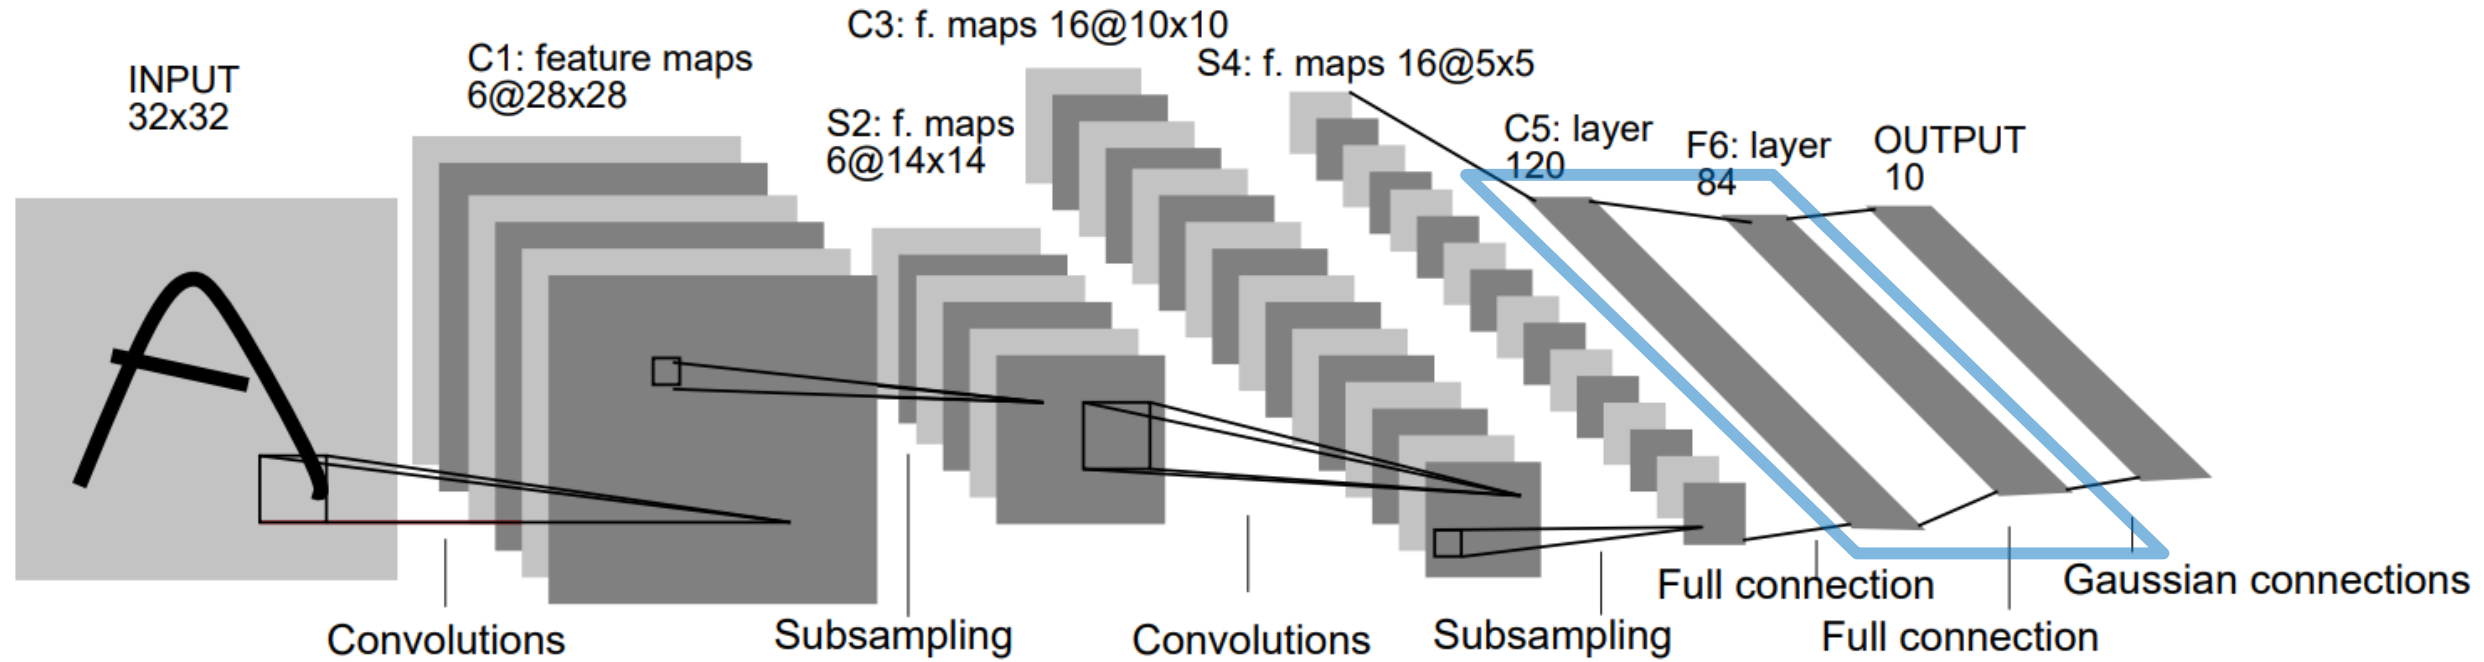


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Step	No. of free parameter	No. of connection
C5→F6 : fully connected	$10,164 = (120+1) * 84$	Same as 'No. of free parameters'

# Case study : LeNet-5

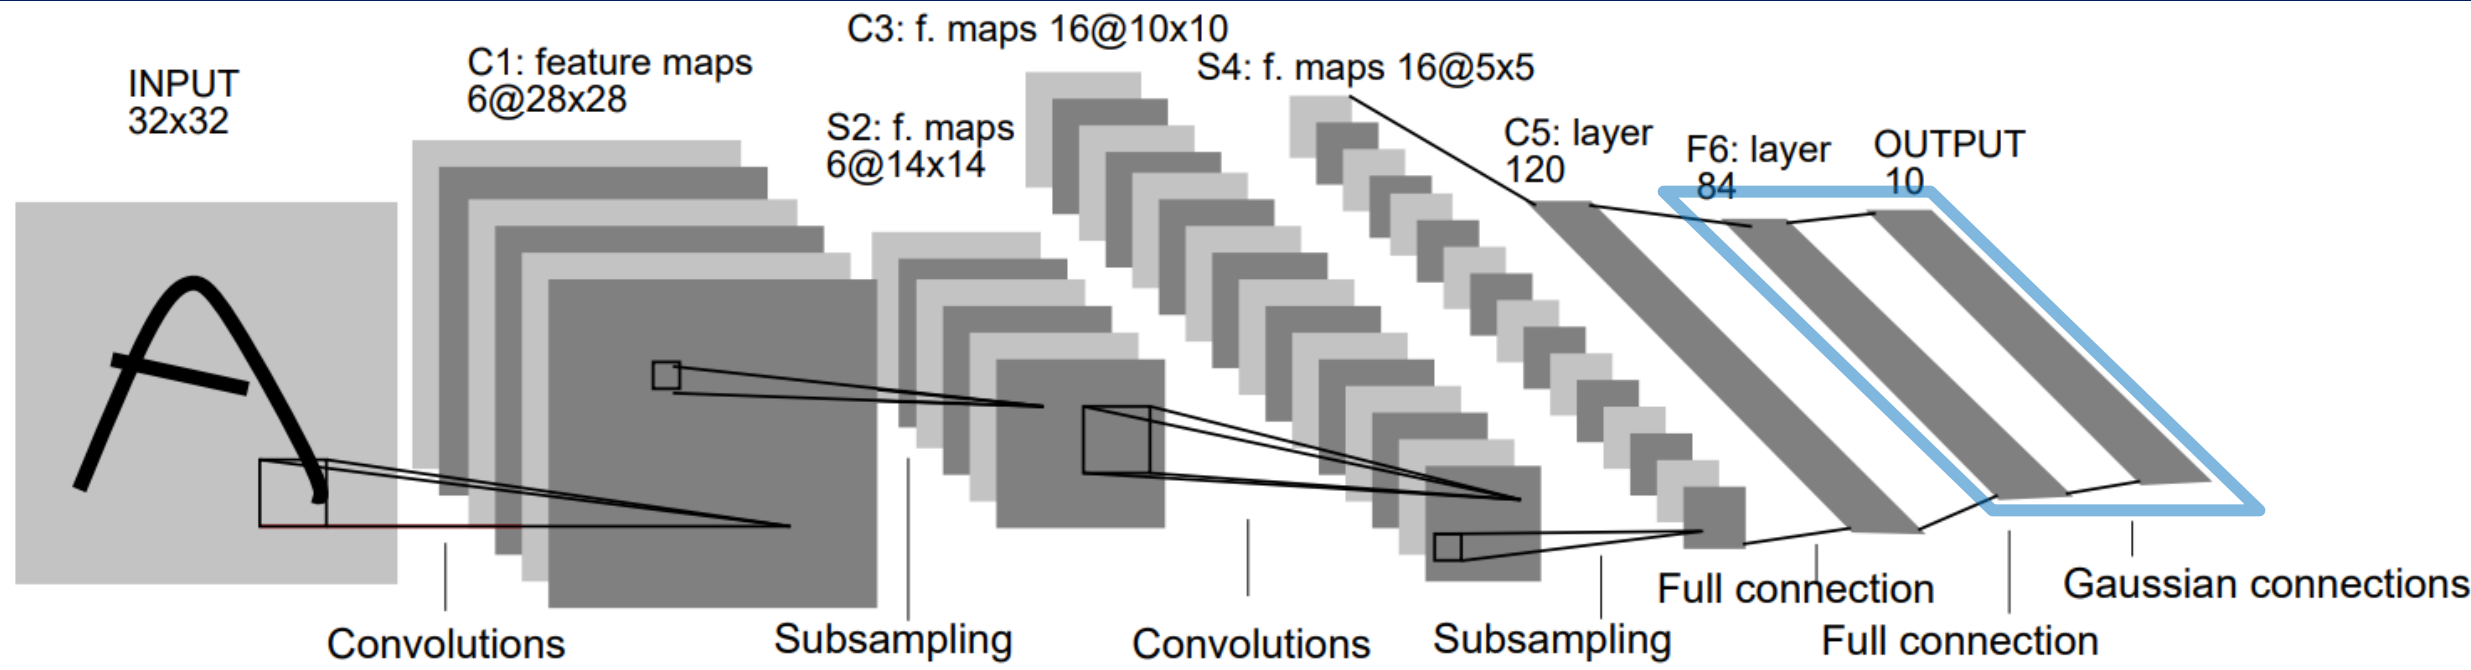


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Step	No. of free parameter	No. of connection
F6→OUTPUT : Gaussian connected	$850=(84+1)*10$	Same as 'No. of free parameters'

# Case study : LeNet-5

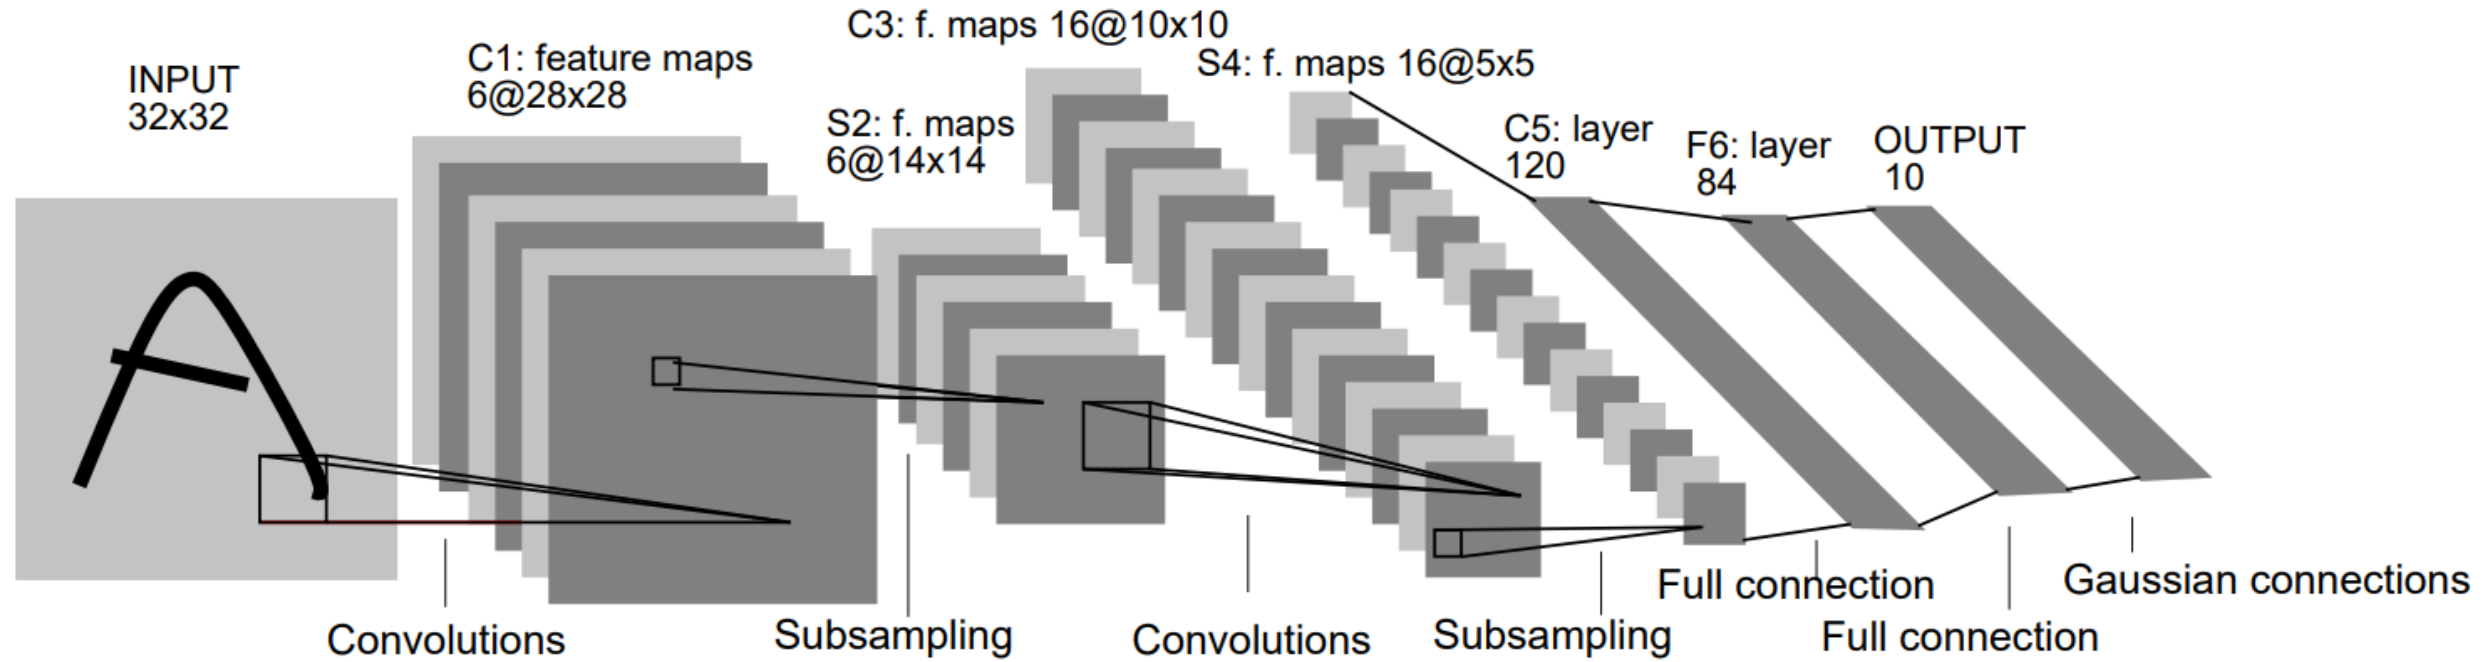


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

To train one image of 32\*32\*1

Total No. of free parameter	Total No. of connection (per image)
60,850	340,918

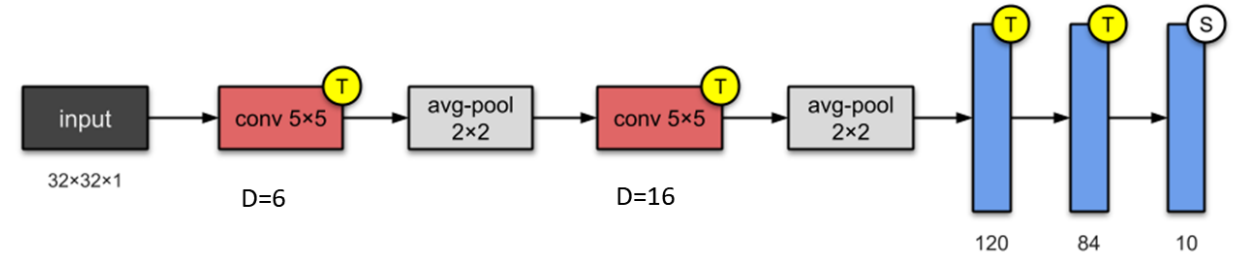
# LeNet-5: Pytorch

- Pytorch implementation

- Go to Tutorial\_LeNet for full description

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        # 1 input channel, 6 output channels, 5x5 convolution kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        # 6*6 from image dimension
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

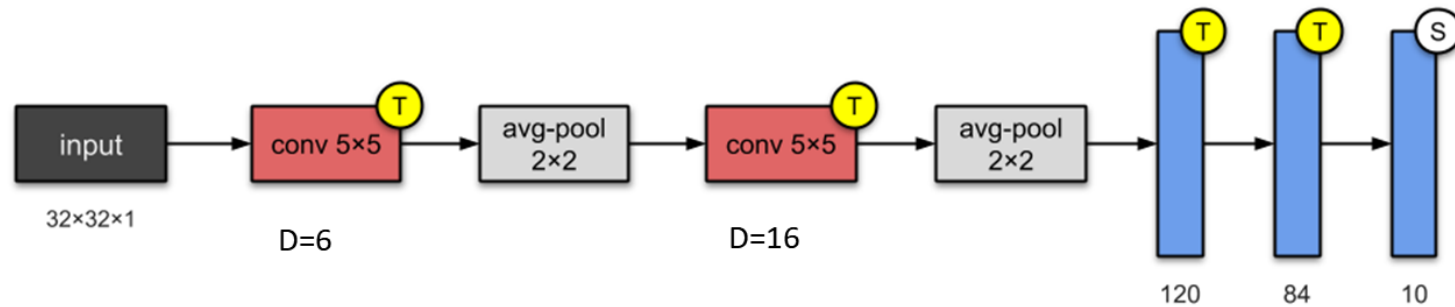
    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```



```
def num_flat_features(self, x):
    size = x.size()[1:]
    # all dimensions except the batch dimension
    num_features = 1
    for s in size:
        num_features *= s
    return num_features
```

# LeNet-5: Keras

- Keras implementation
  - Go to Tutorial\_LeNet for full description



```
model = keras.Sequential()
model.add(layers.Conv2D(filters=6, kernel_size=(5, 5), activation='relu', input_shape=(32,32,1)))
model.add(layers.AveragePooling2D())
model.add(layers.Conv2D(filters=16, kernel_size=(5, 5), activation='relu'))
model.add(layers.AveragePooling2D())
model.add(layers.Flatten())
model.add(layers.Dense(units=120, activation='relu'))
model.add(layers.Dense(units=84, activation='relu'))
model.add(layers.Dense(units=10, activation = 'softmax'))
```

# LeNet-5

## ● Tutorial

- Implement LeNet-5 in PyTorch

Download module: [My\\_DLIP.py](#)

- T2-1: [Create LeNeT CNN model and Train with opendataset \(CIFAR10\)](#)
- T2-2: [Test with loading trained model\(LeNet-5\)](#)

D e e p   L e a r n i n g   I m a g e   P r o c e s s i n g

# Case Study: AlexNet


---




# Case study : AlexNet

## ● AlexNet

- Krizhevsky, Sutskever, Hinton, “Imagenet classification with deep convolutional neural networks”. NIPS 2012
- Winner of ILSVRC 2012
  - Top-5 test error of 15.4% (2<sup>nd</sup> place was 26.2%)
  - A breakthrough in CNN for high classification prediction



ImageNet Top 5 Error Rate



# Case study : AlexNet

## ● AlexNet features

- Multiple GPU
- **ReLU** is introduced
  - instead of TanH
- Local Response Normalization
  - Nowadays use batch normalization
- Overlapping Pooling
- Dropout
- Data augmentation

Input Layer : 224 Width x 224 Height x 3 Channel Image

Convolution 1 : 96 of 11x11x3 Filter + LRN + Pooling

Convolution 2 : 256 of 5x5x48 Filter + LRN + Pooling

Convolution 3 : 384 of 3x3x128 Filter


Convolution 4 : 384 of 3x3x192 Filter

Convolution 5 : 256 of 3x3x192 Filter + Pooling

Fully-Connected 1 : 4096

Fully-Connected 2 : 4096

Output Layer : 1000 Class



Group convolution is applied. From 2nd layer, number of kernels are divided by 2 for each group. e.g.

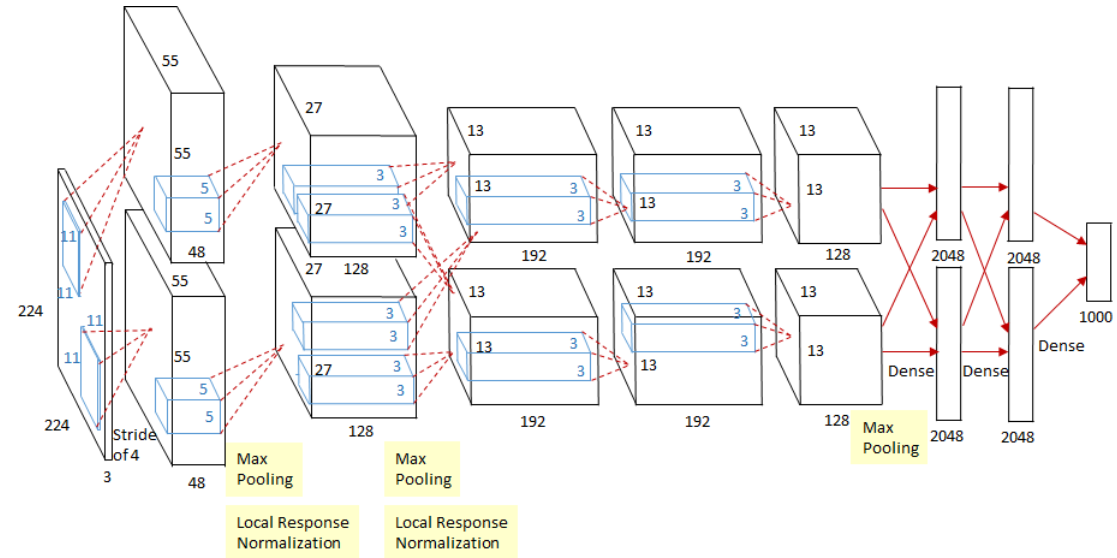
256 of 5x5x48 -->  
(128 of 5x5x48) \* 2

# Case study : AlexNet


## ● AlexNet training

- Dataset
  - ImageNet of 256x256x3
- Input: 224×224×3 image
- Data augmentation
- Initial
  - Bias: 1 for 2,4,5<sup>th</sup> conv. 0 for others
- SGD
  - Batch size: 128
  - Momentum v: 0.9
  - Weight Decay: 0.0005
  - Learning rate : 0.01

Parallel GPU Processings



1 GPU



# Case study : AlexNet

## ● Pytorch implementation

### - Class Alexnet

```
class AlexNet(nn.Module):
```

```
    def __init__(self, num_classes=1000):
        super(AlexNet, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 96, kernel_size=11, stride=4, padding=0),
            nn.ReLU(inplace=True),
            nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(96, 256, kernel_size=5, padding=2, groups=2),
            nn.ReLU(inplace=True),
            nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75),
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(256, 384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(384, 384, kernel_size=3, padding=1, groups=2),
            nn.ReLU(inplace=True),
            nn.Conv2d(384, 256, kernel_size=3, padding=1, groups=2),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
        )
        self.classifier = nn.Sequential(
            nn.Linear(256 * 6 * 6, 4096),
            nn.ReLU(inplace=True),
            nn.Dropout(),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Dropout(),
            nn.Linear(4096, num_classes),
        )
```

```
    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), 256 * 6 * 6)
        x = self.classifier(x)
        return x
```

Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	227x227x3	-	-	-
1	Convolution	96	55 x 55 x 96	11x11	4	relu
	Max Pooling	96	27 x 27 x 96	3x3	2	relu
2	Convolution	256	27 x 27 x 256	5x5	1	relu
	Max Pooling	256	13 x 13 x 256	3x3	2	relu
3	Convolution	384	13 x 13 x 384	3x3	1	relu
4	Convolution	384	13 x 13 x 384	3x3	1	relu
5	Convolution	256	13 x 13 x 256	3x3	1	relu
	Max Pooling	256	6 x 6 x 256	3x3	2	relu
6	FC	-	9216	-	-	relu
7	FC	-	4096	-	-	relu
8	FC	-	4096	-	-	relu
Output	FC	-	1000	-	-	Softmax

# Case study : AlexNet

## ● Keras implementation

- Class Alexnet: see class website

```
model = keras.Sequential()
#1st Convolutional Layer
model.add(Conv2D(filters=96, input_shape=(32,32,3), kernel_size=(11,11), strides=(4,4), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#2nd Convolutional Layer
model.add(Conv2D(filters=256, kernel_size=(5, 5), strides=(1,1), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))

#3rd Convolutional Layer
model.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))

#4th Convolutional Layer
model.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))

#5th Convolutional Layer
model.add(Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
```

	Layer	Feature Map	Size	Kernel Size	Stride	Activation
	Input	Image	1	227x227x3	-	-
1	Convolution	96	55 x 55 x 96	11x11	4	relu
	Max Pooling	96	27 x 27 x 96	3x3	2	relu
2	Convolution	256	27 x 27 x 256	5x5	1	relu
	Max Pooling	256	13 x 13 x 256	3x3	2	relu
3	Convolution	384	13 x 13 x 384	3x3	1	relu
4	Convolution	384	13 x 13 x 384	3x3	1	relu
5	Convolution	256	13 x 13 x 256	3x3	1	relu
	Max Pooling	256	6 x 6 x 256	3x3	2	relu
6	FC	-	9216	-	-	relu
7	FC	-	4096	-	-	relu
8	FC	-	4096	-	-	relu
	Output	FC	-	1000	-	Softmax

```
#Passing it to a Fully Connected layer
model.add(Flatten())
# 1st Fully Connected Layer
model.add(Dense(4096, input_shape=(32,32,3)))
model.add(BatchNormalization())
model.add(Activation('relu'))
# Add Dropout to prevent overfitting
model.add(Dropout(0.4))

#2nd Fully Connected Layer
model.add(Dense(4096))
model.add(BatchNormalization())
model.add(Activation('relu'))
#Add Dropout
model.add(Dropout(0.4))

#3rd Fully Connected Layer
model.add(Dense(1000))
model.add(BatchNormalization())
model.add(Activation('relu'))
#Add Dropout
model.add(Dropout(0.4))

#Output Layer
model.add(Dense(10))
model.add(BatchNormalization())
model.add(Activation('softmax'))

#Model Summary
model.summary()
```